

BEHAVIOUR-ALIGNED FRONT-ENDS FOR FULL-DUPLEX SPEECH LLMs: SIGNAL RESTORATION IS NOT BEHAVIOUR RESTORATION

Jianming Liu

Independent Researcher, Newton, USA

ABSTRACT

Hosted full-duplex speech LLMs are absorbing signal clean-up; what remains on the device is a front-end whose job is to serve the model’s *turn-taking decision*. We ask whether the deployed remedy for noise—a suppressor trained on signal metrics—serves that decision, and what changes when the same 78 k-parameter RNNoise-style front-end is instead aligned to the decision through a frozen model. As the measurement instrument we add a controlled noise protocol to the HumDial test set and score first-response placement. On Lychee-FD (12.8 B), signal-trained suppression restores audibility and with it every clean-condition pathology, while behaviour-level training cuts wrong-place responses from 67 % to 48 % ($p < 10^{-23}$, paired) at no cost in on-time responses. On Freeze-Omni (7B), an eager model whose failures are linguistic, a front-end has no leverage at moderate noise; once masking is severe enough to make noise a trigger, the frozen suppressor makes the model take over earlier and the decision-aligned one later ($p < 10^{-35}$), with first-response relevance preserved. Signal restoration is not behaviour restoration: the two objectives select different front-ends, and the behaviour-aligned one’s leverage tracks the acoustic share of the failure. Aligning from the model’s decisions alone, without gradients, did not reach that solution at the budgets we tried. Protocol, metrics and recipes are released.

Index Terms— full-duplex speech LLM, noise suppression, turn-taking, false barge-in, behaviour-aware front-end

1. INTRODUCTION

Full-duplex spoken dialogue models emit agent audio while ingesting user audio, so their core competence is a continuous decision: speak, keep listening, or back-channel. That decision is made from raw acoustics. For transcription-oriented pipelines, decades of work hardened recognition against noise; the *turn-taking* decision of a full-duplex model has no such history, and it fails differently—not by mis-hearing words but by *acting*: answering a slammed door, treating café babble as a barge-in.

The evidence is already on the record. The ICASSP HumDial challenge [1], whose dual-channel benchmark is the most recent standard for full-duplex interaction, reports that across submitted systems “robustness under background noise remains limited, with transient noises causing false activations or missed detections”—and yet the benchmark itself, like Full-Duplex-Bench and its successors [2, 3, 4], contains no noise condition. The gap reaches the commercial frontier: OpenAI’s GPT-Live-1, released to its API this month, is full-duplex, removes the separate turn detector from the audio path, reports a 30-point gain on Full-Duplex-Bench, and claims “improved handling of ambient noise”—without a single noise-condition number [5]. The community has named the failure and does not yet measure it.

The deployed answer to noise is an upstream suppressor—in the open-source stack, most commonly RNNoise [6] and its descendants, built and tuned against signal metrics on data the downstream model never sees. Hosted full-duplex models change what such a front-end is *for*. Signal clean-up is being absorbed into the model; what stays on the device is what the cloud cannot do: decide, cheaply and privately, whether a sound is addressed to the system and whether a turn is coming—before a metered session sees it. A front-end with that job must be aligned to the *behaviour* of the model it serves. Whether signal-metric training does that, and what behaviour-level alignment changes, are the questions this paper measures, with the model held frozen throughout; whether such alignment is reachable without gradients, as a hosted API would require, we only probe.

A companion line of work on *acoustic echo* found that even cross-entropy on the model’s own turn-taking tokens leaves an objective–behaviour gap dominated by teacher-forcing *exposure*: echo is the model’s own output fed back, and the states that break the benchmark are never visited under teacher forcing (companion work, in submission). Noise carries no such loop—it is exogenous, “do not activate on this” is an ordinary conditional decision fully covered by teacher-forced supervision, and 78 k parameters leave little room for degenerate shortcuts. If behaviour-level training of a front-end works anywhere, it should work here.

Contributions. (1) *Signal restoration is not behaviour restoration*, on two models: the same 78 k front-end, trained on signal metrics, hands a frozen full-duplex model back its clean-

Code, recipes and logs: <https://github.com/jianmliu/behaviour-aligned-frontends>.

condition pathologies (Lychee-FD) or makes it take over earlier (Freeze-Omni, at every SNR tried); aligned to the model’s turn decision through gradients, it reverses that effect, with leverage that tracks the acoustic share of the failure and no cost in first-response relevance. (2) The measurement instrument: a noise-augmented behavioural protocol on the HumDial test set scoring first-response *placement*, grounded in the benchmark’s respond/refrain semantics (§2). (3) On Lychee-FD the behaviour-trained front-end cuts the wrong-place total from 67 to 48 % ($p < 10^{-23}$, paired) at no cost in on-time responses, and ablations show the signal and behaviour solutions to be distinct attractors (§3). (4) Released protocol, metrics, recipes and logs.

2. SETUP

2.1. Model and behaviour semantics

Lychee-FD [7] (12.8 B) emits explicit turn-taking *control tokens* every 0.4 s chunk (keep-listening, start-speaking, backchannel, ...): its behaviour is directly observable and directly supervisable. Audio enters as continuous encoder features, so gradients pass through the frozen model into a front-end.

2.2. Noise protocol on HumDial

The HumDial English test set groups 4,550 recordings into ten scenarios with explicit respond/refrain semantics: four demand a response (ask, repeat, shift, deny); six demand restraint (backchannel, wait, pause, talk-to-others, two others.talk-to-user variants). We take 100/scenario, add noise from held-out DEMAND environments at $\text{SNR} \in [0, 15]$ dB (per-sample seeded: all systems see bit-identical scenes), and feed the frozen model offline. Offline decoding of these long, two-turn recordings yields, in practice, a single model response per sample (99 % of clean decodes), so we measure the *placement* of that response—a quantity with clean semantics per scenario and direct sensitivity to noise: **on-time** (at or after the end of the user’s question), **premature** (inside the question—a failure HumDial itself reports for current models on slow speech), **miss** (no response), **mis-addressed** (placed inside speech directed at someone other than the model: the third-party segments of the *others_** and *talk-to-others* scenarios—the direct false-trigger read-out), and **break-preempt** (inside an annotated mid-question pause). Clean decoding (L0) anchors the distribution; the noise effect is the *paired shift* of this distribution on bit-identical samples, which is robust to the imperfect clean baseline. That baseline is imperfect by design, not by accident: premature take-overs dominate it for both models, which is the behaviour HumDial’s organisers report for current systems on slow or paused speech, and offline decoding of a two-turn recording exposes it fully because nothing interrupts the model once it starts. The metric therefore measures the placement of a




Fig. 1. First-response placement across the four conditions (1000 paired samples). Noise collapses responses into misses; the frozen suppressor restores the clean profile, pathologies included; the behaviour-fine-tuned one shrinks every wrong-place class instead.

decision the models really make, and every comparison in this paper is a within-sample contrast against that same baseline.

2.3. Front-end and training ladder

RN-t is a differentiable RNNNoise-style suppressor: 22 Bark-band features, the original Dense-GRU-GRU-GRU-Dense gain topology (78 k parameters), band gains interpolated to STFT bins, length-preserving. Signal-loss pretraining (SI-SNR, speech from a disjoint corpus, DEMAND ch01) reaches +9.3 dB held-out gain at SNR 5 with 32.5 dB clean-speech fidelity. **R-indep** freezes this checkpoint—deployment practice. **R-joint** fine-tunes it through the frozen model against cross-entropy on the model’s own clean-run control-token trajectory (self-distillation; the released transition-preserving label subsampling), on a training corpus disjoint from HumDial (Full-Duplex-Bench 1.0 user audio, 429 recordings ≤ 20 s, DEMAND noise re-drawn per step). The comparison isolates exactly one variable: what the training objective knows about the model being served.

3. RESULTS

3.1. Noise breaks the turn-taking decision

Table 1, first two columns. The clean baseline already carries the pathologies HumDial reports for current systems—premature answers dominate (59.4 %; 86 % on *talk-to-others*), and 58 % of *others.talk-to-user.before* decodes fire inside speech addressed to a third party. Noise does not add new misbehaviour on top; it *deafens*: misses rise in all ten scenarios (worst on *repeat*, 42 $\rightarrow$ 59), on-time responses halve, and every hearing-dependent failure—premature, mis-addressed, pause preemption—falls passively with the signal.

Table 1. First-response placement on the HumDial test set (1000 samples, 100/scenario; identical noise per sample across all noisy conditions). Mis-addr. = response placed inside third-party speech (false trigger); break-pre. = response inside an annotated mid-question pause (% of the pause scenario). RN-t-frozen = signal-pretrained, frozen (R-indep); RN-t-behav. = same start, fine-tuned on control-token CE (R-joint).

%	clean	noise	+RN-t-frozen	+RN-t-behav.
on-time $\uparrow$	5.4	3.2	6.2	7.0
premature $\downarrow$	59.4	46.4	60.1	43.3
miss $\downarrow$	27.4	44.5	26.6	44.7
mis-addr. $\downarrow$	5.8	4.4	4.8	3.7
break-pre. $\downarrow$	18	10	20	10

The same direction reproduces on VoiceChat-11B under our companion echo protocol’s noise-only control (self-talk 1.60 $\rightarrow$ 1.94, premature take-overs 1.09 $\rightarrow$ 1.39 events/min).

3.2. Frozen versus behaviour-fine-tuned suppression

Table 1, last two columns, one training variable apart. The frozen suppressor is a faithful restorer: misses return to the clean level (26.6 %) and so does everything else—premature back to 60.1 %, pause preemption at 20 %, false triggers at 4.8 %. Signal fidelity hands the model back its clean-condition behaviour, pathologies included. The behaviour-fine-tuned front-end spends the same capacity differently: 17 pt fewer premature answers than the frozen variant (paired exact McNemar on bit-identical samples: 86 vs 252 discordant pairs, $p = 5 \times 10^{-20}$), the lowest false-trigger rate of any condition (3.7 %), and half the pause preemptions—while misses stay at the noise level (44.7 %, $p = 2 \times 10^{-22}$). Aggregating the three wrong-place classes (premature + mis-addressed + break-preempt) gives 66.9 % for the frozen front-end against 48.0 % for the behaviour-trained one ($p = 4 \times 10^{-24}$), at no measurable cost in correct responses: on-time is 7.0 % against the frozen variant’s 6.2 % and the clean baseline’s 5.4 % (both n.s.). The trade it learns is legible: speak when the decision evidence survives the noise, stay silent when it does not, rather than reconstruct everything the model might react to.

3.3. Second model: leverage tracks the acoustic share

To test whether the Lychee-FD finding is specific to one model, we repeat the protocol on Freeze-Omni [8] (7B, frozen Qwen2 backbone, a chunk-level *state head* that decides when to speak; listening needs no token generation, so placement is read from that decision alone) on the same 1,000-sample protocol (200 samples, 20/scenario, for the self-distillation variant and the clean-audio and probe rows), and train the same 78 k front-end from the same signal-pretrained start with a differentiable *expected placement cost* (first-trigger distribution over chunks $\times$




Fig. 2. Freeze-Omni, four systems on 1,000 paired samples per SNR. The signal-trained suppressor makes the eager model take over earlier at every SNR; the behaviour-aligned front-end, same parameters, makes it later, and its gain grows with the acoustic share. Dotted: clean audio.

premature/on-time/miss costs, refrain samples penalising any trigger, plus a logit-margin term in the premature region so that a saturated decision still carries gradient) through the frozen model. Freeze-Omni’s clean profile is the mirror image of Lychee-FD’s: on-time 40.6 %, premature 43.8 %, no misses—and no restraint: it fires inside third-party speech in 99/100 *others_before* samples and inside the pause gap in 57/100. Noise at the protocol’s SNR (0–15 dB) does not move it at all (on-time 40.6 %, premature 43.4 %, restraint axes 100/100 and 60/100): an eager model is not deafened. Pushed further, an acoustic share does appear—but as *more* eagerness, not silence: from –5 to –20 dB premature take-overs rise from 46.3 to 63.0 % (Fig. 2; 1,000 samples each) with still no misses; sustained masking acts as a trigger, whereas 300 ms bursts at +10/+20 dB barely register (42.5/40.0, 39.5/43.5): transient false activations are not this model’s failure mode. Two consequences follow. Self-distillation cannot be the target here—the model’s own clean first response precedes the annotated turn end by a median 7.3 s on Full-Duplex-Bench—so we take the benchmark’s turn-end annotations as placement targets; and the front-end’s job is restraint rather than restoration, so we also check clean audio, where all three front-ends leave placement unchanged (200 samples, n.s.). Under noise, the frozen suppressor nudges the model slightly *more* eager (on-time 38.8 %, premature 45.1 %, restraint axes 100/100 and 61/100: cleaner speech, if anything earlier take-over). The behaviour-trained front-ends land on the noisy baseline—soft placement with annotation targets 39.7/44.3 on 1,000 samples, restraint axes 100/100 and 60/100 exactly as without it; on 200 samples, self-distillation targets 43.0/39.0 (on-time/premature), all within sampling noise. Per sample, the front-ends move decision instants the *wrong* way for restraint—earlier take-overs outnumber later ones 172:115 (frozen) and 156:91 (trained), sign test $p < 10^{-3}$: cleaner speech makes an eager model marginally more eager. The trained suppressor kept 6.4 of the frozen model’s 8.6 dB SI-SNR gain on its training mixtures: nothing collapsed, and the only direction found was the acous-

tic one. *Where the acoustic share exists, the leverage appears.* At -15 dB (1,000 paired samples; 28.4/56.5 without a front-end) the frozen suppressor makes matters worse—24.3/61.2, decisions earlier 340:174 (sign test $p = 2 \times 10^{-13}$): denoising feeds the trigger. The behaviour-trained front-end moves them the other way: 35.8/49.0, later 309:173 against no front-end ($p = 6 \times 10^{-10}$) and 432:139 against the frozen suppressor ($p < 10^{-35}$). The same 78 k parameters, aligned to the decision instead of the signal, reverse the front-end’s effect: the Lychee-FD finding on a second model. A hosted API would expose only the decisions; as a probe, evolution strategies on the realised placement reward (23 band gains, 3k rollouts) end at the frozen solution (24.7/60.8, $p = 0.27$ vs. frozen), and neither tripling the budget, searching all 78 k parameters, nor using the gradient arm’s own loss as reward moves them (24.6/60.5, 17.4/68.7, 25.9/59.5): at these budgets, decision-only feedback did not find the direction gradients find. The pattern replicates at -10 and -20 dB (all sign tests $p < 10^{-6}$); at -5 dB, where noise alone barely moves the model, gradient alignment only cancels the frozen suppressor’s harm (Fig. 2). Content follows: judged by Qwen2-7B-Instruct on the opening of the first response (0–2 relevance, -15 dB), the frozen suppressor lowers relevance (0.77 vs. 0.83, $p = 0.07$), gradient alignment keeps it (0.84; 170:132 above frozen, $p = 0.03$) and cuts hallucinated Chinese openings from 13.6 to 9.4 % ($p = 2 \times 10^{-4}$): later take-overs cost no relevance. Across regimes, a waveform front-end has leverage only over the acoustic share of a failure—large on Lychee-FD, near zero here.

3.4. Ablations: two attractors, no middle ground

Three variants probe what the behaviour objective needs. (i) *Gentler learning rate* (3×10^{-5} vs 10^{-4}): the fine-tune never leaves the pretrained signal solution—its profile is indistinguishable from the frozen front-end (on-time 6.2 %, premature 62 %, miss 25 %). (ii) *Scratch start*: without signal pretraining the same objective learns nothing usable; the profile matches raw noise (on-time 3.6 % vs 3.2 %). (iii) *SI-SNR anchor* added to the CE ($\lambda=0.05$): the anchor dominates and the solution reverts to the frozen profile, restraint gains included. The picture is consistent: the signal solution and the behaviour solution are distinct attractors—behaviour training must start from signal competence (ii), needs enough learning rate to escape it (i), and is not reachable as a regularised compromise (iii). The objective is not a seasoning over signal fidelity; it selects a different front-end.

4. DISCUSSION

Where front-ends still pay. Egocentric-device challenges chart the decline of the generic front-end: in CHiME-8 MM-CSG [9] training data beat array processing on WER, yet a neural separation front-end won the *speaker-attribution* axis [10];

by SLT 2026 SmartGlasses [11] the winner used no front-end at all. The opposite reaction grows the front-end into a 30B LLM trained on task labels (UAF [12]). Our results locate the surviving value: a front-end pays on the axis its objective shares with the downstream decision—attribution there, turn-taking here—and 78 k parameters aligned to that decision shift behaviour where signal restoration does not. Deployment follows: with hosted full-duplex models metered per minute of audio, every noise-induced wrong-place response is a billed event, and an on-device front-end that removes them at zero marginal cost is a cost control, not an audio-quality option. The same judgement—is this addressed to the system, is a turn coming—is what an on-device gate needs to decide when a metered session should be open at all; we leave that gate to future work. Unlike enhancement-for-ASR joint training, the model is frozen, the target is behaviour rather than transcription, and the failure repaired (acting at the wrong time) has no transcription analogue.

Relation to the echo domain. In the companion echo setting the same CE failed to move free-running behaviour because teacher forcing never visits the states echo creates; here the perturbation is exogenous and the same objective moves behaviour—consistent with the exposure account. What remains is an *objective* gap: the CE optimum under noise is a restraint profile, reachable only from a signal-competent start (§3.4).

Limitations. English subset only; offline decoding of the official test set (no streaming latency axis); front-ends trained through two models (the VoiceChat corroboration is measurement-only); single noise corpus (DEMAND), though held out between pretraining and test; one LLM judge, opening sentences only; the gradient-free probe is zeroth-order search at small budgets on one model.

5. CONCLUSION

We measured what noise does to the turn-taking decision of a frozen full-duplex speech LLM on the benchmark that first named the problem, and compared the deployed remedy—a signal-trained suppressor—against the same suppressor fine-tuned on the model’s own turn-taking tokens. Signal training restores what the model hears; behaviour training changes what the model does: the wrong-place total falls from 67 to 48 % ($p < 10^{-23}$, paired), on-time responses are preserved, and the cost is silence where the evidence does not survive the noise. The two solutions are separate attractors, not ends of a dial. On a second, eager model the same alignment reverses the frozen suppressor’s effect once masking creates an acoustic failure. This reframes the front-end question for deployment: the choice is not how much to denoise, but which decision the denoiser should serve.

Use of AI. Claude (Anthropic) assisted with code, experiment orchestration and editing; every result and claim was checked by the author, who takes full responsibility for them.

6. REFERENCES

- [1] Chengyou Wang, Hongfei Xue, Guojian Li, Zhixian Zhao, Shuiyuan Wang, Shuai Wang, Xin Xu, Hui Bu, and Lei Xie, “Full-duplex interaction in spoken dialogue systems: A comprehensive study from the ICASSP 2026 HumDial challenge,” *arXiv preprint arXiv:2604.21406*, 2026.
- [2] Guan-Ting Lin, Jiachen Lian, Tingle Li, Qirui Wang, Gopala Anumanchipalli, Alexander H. Liu, and Hung-yi Lee, “Full-duplex-bench: A benchmark to evaluate full-duplex spoken dialogue models on turn-taking capabilities,” *arXiv preprint arXiv:2503.04721*, 2025.
- [3] Guan-Ting Lin, Shih-Yun Shan Kuan, Qirui Wang, Jiachen Lian, Tingle Li, and Hung-yi Lee, “Full-duplex-bench v1.5: Evaluating overlap handling for full-duplex speech models,” *arXiv preprint arXiv:2507.23159*, 2025.
- [4] Guan-Ting Lin, Chen Chen, Zhehuai Chen, and Hung-yi Lee, “Full-duplex-bench-v3: Benchmarking tool use for full-duplex voice agents under real-world disfluency,” *arXiv preprint arXiv:2604.04847*, 2026.
- [5] OpenAI, “Build more natural voice experiences with GPT-Live-1 in the API,” <https://openai.com/index/introducing-gpt-live-1-in-the-api/>, Sept. 2026, Accessed 2026-09-11.
- [6] Jean-Marc Valin, “A hybrid DSP/deep learning approach to real-time full-band speech enhancement,” in *Proc. IEEE MMSP*, 2018.
- [7] Zhenyu Liu, Xuanyu Zhang, Yunxin Li, et al., “Lychee-fd: Hierarchical acoustic-semantic modeling for full-duplex spoken language models,” <https://arxiv.org/abs/2607.06540>, 2026, Weights: huggingface.co/HIT-TMG/Lychee-FD.
- [8] Xiong Wang, Yangze Li, Chaoyou Fu, Yunhang Shen, Lei Xie, Ke Li, Xing Sun, and Long Ma, “Freeze-omni: A smart and low latency speech-to-speech dialogue model with frozen LLM,” *arXiv preprint arXiv:2411.00774*, 2024.
- [9] Katerina Zmolikova, Simone Merello, Kaustubh Kalganekar, Ju Lin, Niko Moritz, Pingchuan Ma, Ming Sun, Honglie Chen, Antoine Saliou, Stavros Petridis, Christian Fuegen, and Michael Mandel, “The CHiME-8 MMCSG challenge: Multi-modal conversations in smart glasses,” in *Proc. 8th Int. Workshop on Speech Processing in Everyday Environments (CHiME)*, 2024.
- [10] Chenda Pang, Feifei Xiong, Yuxin Ni, Lin Zhou, and Jinwei Feng, “The SEUEE system for the CHiME-8 MMCSG challenge,” in *Proc. 8th Int. Workshop on Speech Processing in Everyday Environments (CHiME)*, 2024.
- [11] Dehui Gao, Zhixian Zhao, Zhennan Lin, Yujie Liao, Yuhang Dai, Yike Zhu, Longshuai Xiao, Hui Bu, Xin Xu, Xie Chen, Shuai Wang, Liumeng Xue, Zhonghua Fu, Jun Du, Eng-Siong Chng, Jun Zhou, and Lei Xie, “The SLT 2026 SmartGlasses challenge: Benchmarking egocentric multi-talker speech recognition and understanding with audio-language models,” *arXiv preprint arXiv:2608.12034*, 2026.
- [12] Yadong Li, Guoxin Wu, Haiping Hou, and Biye Li, “UAF: A unified audio front-end LLM for full-duplex speech interaction,” *arXiv preprint arXiv:2604.19221*, 2026.